

Checkpoint 4

December 1, 2023

1 Project title: Health Tales

1.0.1 Team members

1. Adit Doshi
2. Yash Kharade
3. Harshaditya Mallipudi
4. Prathik Makthala

The dataset we are using is collected from National Health Insurance Service in Korea. It has in total 23 columns.

1. sex - male, female
2. age - rounded to 5 years
3. height - rounded upto 5 cms
4. weight - in kilograms
5. sight_left - left eye sight
6. sight_right - right eye sight
7. hear_left - hearing left, 1(normal), 2(abnormal)
8. hear_right - hearing right, 1(normal), 2(abnormal)
9. SBP - Systolic blood pressure
10. DBP - Diastolic blood pressure
11. BLDS - BLDS or FSG(fasting blood glucose)
12. tot_chole - total cholesterol
13. HDL_chole - HDL cholesterol
14. LDL_chole - LDL cholesterol
15. triglyceride
16. hemoglobin
17. urine_protein - protein in urine, 1(-), 2(+/-), 3(+1), 4(+2), 5(+3), 6(+4)
18. serum_creatinine - serum(blood) creatinine
19. SGOT_AST - SGOT(Glutamate-oxaloacetate transaminase) AST(Aspartate transaminase)
20. SGOT_ALT - ALT(Alanine transaminase)
21. gamma_GTP - γ -glutamyl transpeptidase
22. SMK_stat_type_cd - Smoking state, 1(never), 2(used to smoke but quit), 3(still smoke)
23. DRK_YN - Drinker or Not

1.1 Importing Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import classification_report, accuracy_score, \
    roc_auc_score, roc_curve
```

1.2 Loading data from CSV

```
[2]: df = pd.read_csv("smoking_driking_dataset_Ver01.csv")
```

1.3 Data Cleaning

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 991346 entries, 0 to 991345
Data columns (total 24 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   sex                   991346 non-null object
 1   age                   991346 non-null int64
 2   height                991346 non-null int64
 3   weight                991346 non-null int64
 4   waistline             991346 non-null float64
 5   sight_left            991346 non-null float64
 6   sight_right           991346 non-null float64
 7   hear_left             991346 non-null float64
 8   hear_right            991346 non-null float64
 9   SBP                   991346 non-null float64
10  DBP                   991346 non-null float64
11  BLDS                  991346 non-null float64
12  tot_chole              991346 non-null float64
13  HDL_chole              991346 non-null float64
14  LDL_chole              991346 non-null float64
15  triglyceride           991346 non-null float64
16  hemoglobin             991346 non-null float64
17  urine_protein          991346 non-null float64
18  serum_creatinine       991346 non-null float64
19  SGOT_AST               991346 non-null float64
```

```

20  SGOT_ALT          991346 non-null  float64
21  gamma_GTP         991346 non-null  float64
22  SMK_stat_type_cd  991346 non-null  float64
23  DRK_YN            991346 non-null  object
dtypes: float64(19), int64(3), object(2)
memory usage: 181.5+ MB

```

Apart from sex and the DRN_YK(target variable) feature all of the features are quantative.

There is no null value in any of the columns.

```
[4]: df.nunique()
```

```

[4]: sex                2
     age                14
     height             13
     weight             24
     waistline          737
     sight_left         24
     sight_right        24
     hear_left          2
     hear_right         2
     SBP                171
     DBP                127
     BLDS               498
     tot_chole          474
     HDL_chole          223
     LDL_chole          432
     triglyceride       1657
     hemoglobin         190
     urine_protein       6
     serum_creatinine   183
     SGOT_AST           568
     SGOT_ALT           594
     gamma_GTP          940
     SMK_stat_type_cd    3
     DRK_YN             2
     dtype: int64

```

```
[5]: df.drop_duplicates(inplace=True)
```

```
[6]: df.dropna(inplace=True)
```

```
[7]: df_for_eda = pd.get_dummies(df, drop_first=True)
```

1.4 Exploratory Data Analysis

Below are the min, max, mean and other values for all the features:

```
[8]: df_for_eda.describe().transpose()
```

```
[8]:
```

	count	mean	std	min	25%	50%	
age	991320.0	47.614529	14.181346	20.0	35.0	45.0	\
height	991320.0	162.240563	9.282922	130.0	155.0	160.0	
weight	991320.0	63.283884	12.514101	25.0	55.0	60.0	
waistline	991320.0	81.233255	11.850296	8.0	74.1	81.0	
sight_left	991320.0	0.980833	0.605954	0.1	0.7	1.0	
sight_right	991320.0	0.978428	0.604779	0.1	0.7	1.0	
hear_left	991320.0	1.031495	0.174652	1.0	1.0	1.0	
hear_right	991320.0	1.030476	0.171892	1.0	1.0	1.0	
SBP	991320.0	122.432360	14.543083	67.0	112.0	120.0	
DBP	991320.0	76.052549	9.889334	32.0	70.0	76.0	
BLDS	991320.0	100.424305	24.179852	25.0	88.0	96.0	
tot_chole	991320.0	195.556769	38.660092	30.0	169.0	193.0	
HDL_chole	991320.0	56.936984	17.238578	1.0	46.0	55.0	
LDL_chole	991320.0	113.037429	35.842938	1.0	89.0	111.0	
triglyceride	991320.0	132.140030	102.194762	1.0	73.0	106.0	
hemoglobin	991320.0	14.229810	1.584924	1.0	13.2	14.3	
urine_protein	991320.0	1.094221	0.437719	1.0	1.0	1.0	
serum_creatinine	991320.0	0.860467	0.480536	0.1	0.7	0.8	
SGOT_AST	991320.0	25.989424	23.493668	1.0	19.0	23.0	
SGOT_ALT	991320.0	25.755148	26.308910	1.0	15.0	20.0	
gamma_GTP	991320.0	37.136152	50.423811	1.0	16.0	23.0	
SMK_stat_type_cd	991320.0	1.608112	0.818504	1.0	1.0	1.0	

	75%	max
age	60.0	85.0
height	170.0	190.0
weight	70.0	140.0
waistline	87.8	999.0
sight_left	1.2	9.9
sight_right	1.2	9.9
hear_left	1.0	2.0
hear_right	1.0	2.0
SBP	131.0	273.0
DBP	82.0	185.0
BLDS	105.0	852.0
tot_chole	219.0	2344.0
HDL_chole	66.0	8110.0
LDL_chole	135.0	5119.0
triglyceride	159.0	9490.0
hemoglobin	15.4	25.0
urine_protein	1.0	6.0
serum_creatinine	1.0	98.0
SGOT_AST	28.0	9999.0
SGOT_ALT	29.0	7210.0

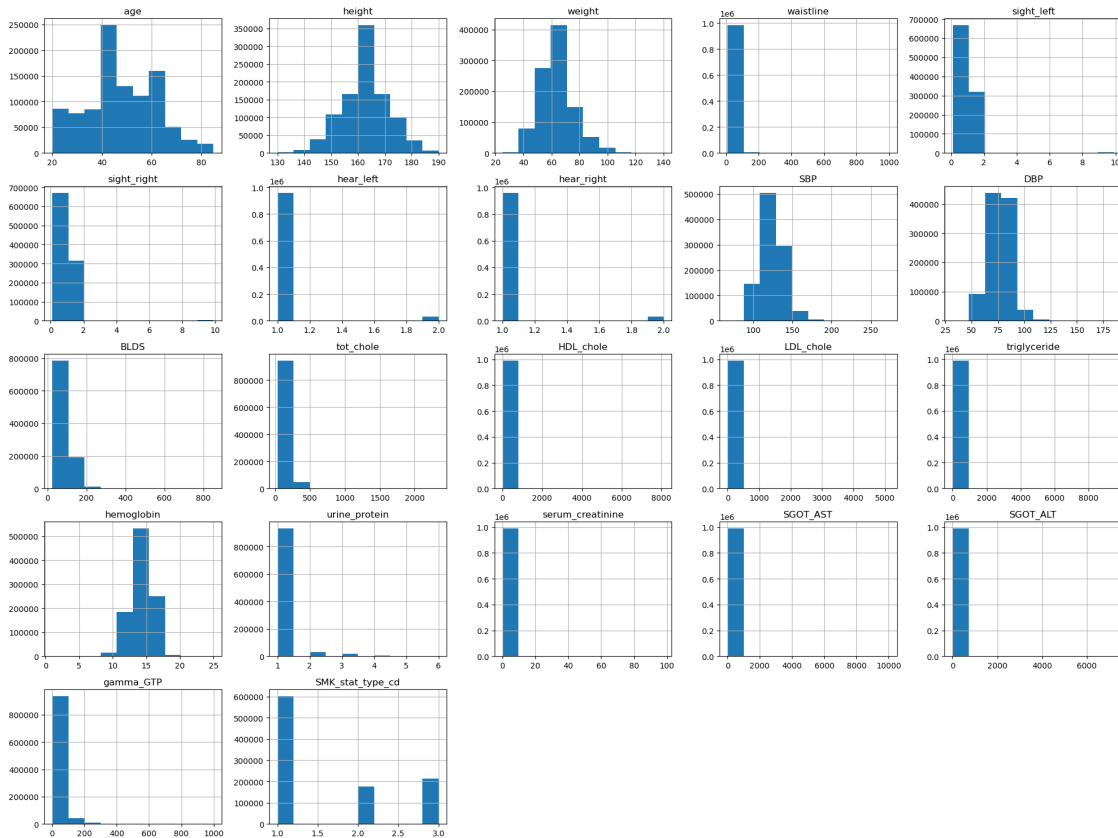
gamma_GTP	39.0	999.0
SMK_stat_type_cd	2.0	3.0

Observations:

1. waistline: The maximum value of 999 cm in the waistline column appears to be an outlier or potentially an error, as it is significantly higher than the other values.
2. HDL_chole: The extremely high maximum value of 8110 mg/dL for HDL cholesterol seems unusual and could be an outlier.
3. serum_creatinine: The maximum value of 98 mg/dL is considerably higher than the other values in the serum creatinine column.
4. SGOT_AST and SGOT_ALT: Both SGOT and SGOT levels have very high maximum values (9999 and 7210, respectively), which could be outliers or data entry errors.
5. gamma_GTP: The maximum value of 999 is significantly higher than the other values, indicating a potential outlier.
6. BLDS: The maximum value of 852 for blood sugar appears to be significantly higher than the typical range.

Histogram showing frequency of all the features:

```
[9]: df_for_eda.hist(figsize=(20,15))  
plt.tight_layout()  
plt.show()
```



```
[10]: df_for_eda.skew()
```

```
[10]: age                0.153652
      height            -0.022721
      weight            0.576537
      waistline         26.789317
      sight_left        9.994631
      sight_right       10.033672
      hear_left         5.365007
      hear_right        5.463034
      SBP               0.482035
      DBP               0.400013
      BLDs              4.617495
      tot_chole         1.556927
      HDL_chole        104.578552
      LDL_chole         5.251830
      triglyceride      6.529776
      hemoglobin        -0.383984
      urine_protein      5.672664
      serum_creatinine  111.021200
```

```

SGOT_AST          150.490203
SGOT_ALT          50.038404
gamma_GTP         7.718657
SMK_stat_type_cd  0.831452
sex_Male          -0.124272
DRK_YN_Y          0.000742
dtype: float64

```

Symmetric Distributions (Skewness 0): age: The distribution of age appears to be approximately symmetric with a skewness close to 0, indicating that it's balanced around the mean.

height: Similarly, the height distribution is very close to symmetric.

Right-Skewed Distributions (Positive Skewness): weight, SBP (Systolic Blood Pressure), DBP (Diastolic Blood Pressure), tot_chole (Total Cholesterol), LDL_chole (LDL Cholesterol), triglyceride, hemoglobin: These columns exhibit right-skewed distributions, as indicated by positive skewness values. This means that the majority of data points are concentrated on the left side of the distribution, with a long tail on the right.


Strongly Right-Skewed Distributions: waistline, sight_left, sight_right, hear_left, hear_right, BLDS (Blood Sugar), HDL_chole (HDL Cholesterol), serum_creatinine, SGOT_AST (Aspartate Aminotransferase), gamma_GTP: These columns have very high positive skewness values, indicating that they are strongly right-skewed. These variables have long right tails with most values concentrated on the lower end.

Left-Skewed Distribution (Negative Skewness): sex_Male: This column has a slightly negative skewness, suggesting that there may be more females than males in the dataset. A negative skew means that the data is concentrated on the right side with a long tail on the left.

```
[11]: df_for_eda['BMI'] = df_for_eda['weight'] / (df_for_eda['height']/100)**2
```

```
[12]: # Visualize correlation matrix
corr_matrix = df_for_eda.corr()
mask = np.triu(np.ones_like(corr_matrix, dtype=bool))

plt.figure(figsize=(17, 17))
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", mask = mask)
plt.show()
```



None of the features have a strong coorelation with the target variable

Splitting the data:

```
[13]: X = df.drop("DRK_YN", axis=1)

y = df["DRK_YN"].map({'Y': 1, 'N': 0})

X_encoded = pd.get_dummies(X, drop_first=True)

df_encoded = pd.concat([X_encoded, y], axis=1)
```



```

# Computing the correlation matrix
corr_matrix = df_encoded.corr()

correlations_with_target = corr_matrix["DRK_YN"].drop("DRK_YN")

threshold = 0.2
selected_features = correlations_with_target[correlations_with_target.abs() >=
↳threshold].index.tolist()

X_selected = X_encoded[selected_features]

# Split data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X_selected, y, test_size=0.
↳3, random_state=42)

# Normalize data
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

```

1.4.1 Model Training

For preliminary analysis, we have selected Logistic Regression, Random forest and K-nearest neighbours models for our data.

Below are some results and plotted ROC curves for all three algorithms

```

[31]: from sklearn.metrics import accuracy_score, roc_auc_score, roc_curve,
↳classification_report
import matplotlib.pyplot as plt

models = {
    "Logistic Regression": LogisticRegression(),
    "Random Forest": RandomForestClassifier(),
    "K-Nearest Neighbors": KNeighborsClassifier()
}

original_model_accuracies = {}

for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred_test = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred_test)
    original_model_accuracies[name] = accuracy # Store each model's accuracy
↳in the dictionary

    print(f"Model: {name}")
    print(classification_report(y_test, y_pred_test))

```

```

print("The TEST accuracy is", accuracy)

# Compute and display ROC AUC score
roc_score = roc_auc_score(y_test, y_pred_test)
print("The ROC score for TEST data is", roc_score)

# Plot ROC curve
fpr, tpr, thresholds = roc_curve(y_test, model.predict_proba(X_test)[:,-1])
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label=f'ROC AUC = {roc_score:.2f}')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(f'ROC Curve - {name}')
plt.legend(loc='lower right')
plt.grid()
plt.show()
if name == "Random Forest":
    # Using feature_importances_ for RandomForestClassifier
    feature_names = X_selected.columns
    plt.figure(figsize=(10, 8))
    plt.barh(feature_names, model.feature_importances_)
    plt.xlabel('Importance')
    plt.ylabel('Feature')
    plt.title(f'Feature Importances for {name}')
    plt.show()


```

Model: Logistic Regression

	precision	recall	f1-score	support
0	0.71	0.72	0.72	148796
1	0.72	0.70	0.71	148600
accuracy			0.71	297396
macro avg	0.71	0.71	0.71	297396
weighted avg	0.71	0.71	0.71	297396

The TEST accuracy is 0.7129416670029187

The ROC score for TEST data is 0.7129337993246248





Model: Random Forest

	precision	recall	f1-score	support
0	0.68	0.67	0.68	148796
1	0.67	0.69	0.68	148600
accuracy			0.68	297396
macro avg	0.68	0.68	0.68	297396
weighted avg	0.68	0.68	0.68	297396

The TEST accuracy is 0.6786607755316144

The ROC score for TEST data is 0.6786680265938638



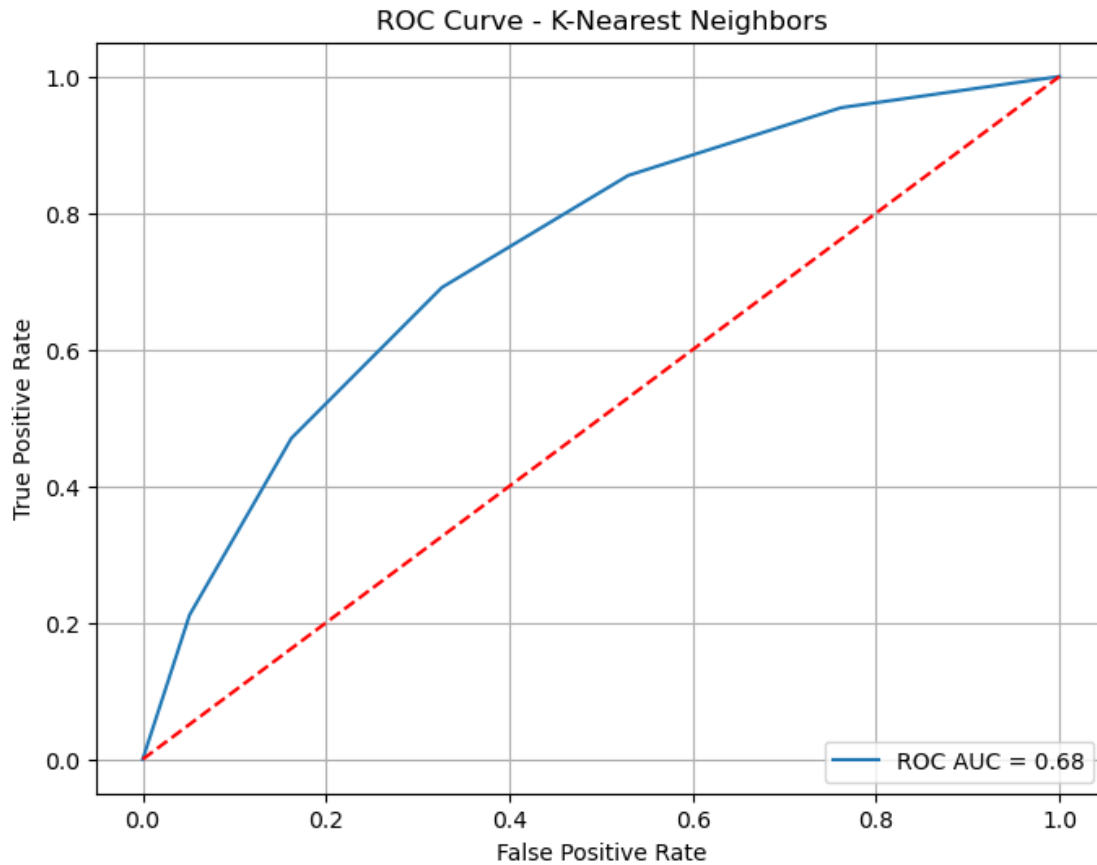


Model: K-Nearest Neighbors

	precision	recall	f1-score	support
0	0.69	0.67	0.68	148796
1	0.68	0.69	0.68	148600
accuracy			0.68	297396
macro avg	0.68	0.68	0.68	297396
weighted avg	0.68	0.68	0.68	297396


The TEST accuracy is 0.682315834779217

The ROC score for TEST data is 0.682321551683559



Accuracy comparison for all the three models. Without any fine tuning, Logistic regression gave us the best accuracy followed by KNN and random forest.

```
[15]: plt.figure(figsize=(10, 6))
plt.bar(models.keys(), model_accuracies, color=['blue', 'green', 'red'])
plt.xlabel('Models')
plt.ylabel('Accuracy')
plt.title('Model Accuracy Comparison')
plt.ylim([min(model_accuracies) - 0.05, max(model_accuracies) + 0.05])
plt.show()
```



General Observations:

1. The Logistic Regression model has the highest accuracy, followed by K-Nearest Neighbors and Random Forest.
2. Precision, recall, and F1-scores are similar for all three models, suggesting that they perform similarly in terms of correctly classifying instances from both classes.
3. The ROC scores for all models are also similar, indicating comparable discrimination abilities.

1.5 Set Up Logistic Regression and GridSearchCV

```
[22]: import warnings
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler

warnings.filterwarnings('ignore')

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

C_values = [0.001, 0.01, 0.1, 1, 10, 100]
solvers = ['lbfgs', 'newton-cg', 'sag', 'saga']
max_iters = [100, 200, 300, 500] # Increased max_iter
```

```

results = []

for C in C_values:
    for solver in solvers:
        for max_iter in max_iters:
            try:
                logreg = LogisticRegression(C=C, solver=solver,
↪max_iter=max_iter)
                logreg.fit(X_train_scaled, y_train)
                y_pred = logreg.predict(X_test_scaled)
                accuracy = accuracy_score(y_test, y_pred)

                results.append({'C': C, 'solver': solver, 'max_iter': max_iter,
↪'accuracy': accuracy})
            except ValueError:
                continue

sorted_results = sorted(results, key=lambda x: x['accuracy'], reverse=True)

for result in sorted_results[:5]:
    print(result)

```

```

{'C': 1, 'solver': 'sag', 'max_iter': 200, 'accuracy': 0.7129618421229607}
{'C': 10, 'solver': 'sag', 'max_iter': 200, 'accuracy': 0.7129584796029537}
{'C': 100, 'solver': 'sag', 'max_iter': 200, 'accuracy': 0.7129584796029537}
{'C': 0.1, 'solver': 'sag', 'max_iter': 100, 'accuracy': 0.7129551170829467}
{'C': 1, 'solver': 'saga', 'max_iter': 200, 'accuracy': 0.7129551170829467}

```

```

[21]: import warnings
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score

warnings.filterwarnings('ignore')

param_grid = {
    'max_depth': [None, 10, 20],
    'max_features': ['auto'],
    'min_samples_leaf': [1, 2],
    'min_samples_split': [2, 5],
    'n_estimators': [10, 50, 100]
}

rf = RandomForestClassifier()
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5,
↪n_jobs=-1, verbose=2)

```



```

grid_search.fit(X_train, y_train)
print("Best Parameters:", grid_search.best_params_)

best_rf = grid_search.best_estimator_
y_pred = best_rf.predict(X_test)

best_accuracy = accuracy_score(y_test, y_pred)
print("Accuracy of the best model:", best_accuracy)

print("\nAccuracies for all parameter combinations:")
for i, params in enumerate(grid_search.cv_results_['params']):
    print(f"Parameters: {params} - Accuracy: {grid_search.
cv_results_['mean_test_score'][i]}")

```

Fitting 5 folds for each of 36 candidates, totalling 180 fits
Best Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}
Accuracy of the best model: 0.721761556981264

Accuracies for all parameter combinations:

```

Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 1,
'min_samples_split': 2, 'n_estimators': 10} - Accuracy: 0.667940582308622
Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 1,
'min_samples_split': 2, 'n_estimators': 50} - Accuracy: 0.6770727067508013
Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 1,
'min_samples_split': 2, 'n_estimators': 100} - Accuracy: 0.6784402921222288
Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 1,
'min_samples_split': 5, 'n_estimators': 10} - Accuracy: 0.6843804218137597
Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 1,
'min_samples_split': 5, 'n_estimators': 50} - Accuracy: 0.6918653918394967
Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 1,
'min_samples_split': 5, 'n_estimators': 100} - Accuracy: 0.6929130577866849
Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 2,
'min_samples_split': 2, 'n_estimators': 10} - Accuracy: 0.691052624823594
Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 2,
'min_samples_split': 2, 'n_estimators': 50} - Accuracy: 0.6976859139425777
Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 2,
'min_samples_split': 2, 'n_estimators': 100} - Accuracy: 0.6987753701424188
Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 2,
'min_samples_split': 5, 'n_estimators': 10} - Accuracy: 0.6933021489202353
Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 2,
'min_samples_split': 5, 'n_estimators': 50} - Accuracy: 0.6995391427516857
Parameters: {'max_depth': None, 'max_features': 'auto', 'min_samples_leaf': 2,
'min_samples_split': 5, 'n_estimators': 100} - Accuracy: 0.7001905110023451
Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 1,
'min_samples_split': 2, 'n_estimators': 10} - Accuracy: 0.7203872473795088
Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 1,

```

'min_samples_split': 2, 'n_estimators': 50} - Accuracy: 0.7211827236200854
 Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 1,
 'min_samples_split': 2, 'n_estimators': 100} - Accuracy: 0.7214867915194109
 Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 1,
 'min_samples_split': 5, 'n_estimators': 10} - Accuracy: 0.720529914395336
 Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 1,
 'min_samples_split': 5, 'n_estimators': 50} - Accuracy: 0.7211711956716093
 Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 1,
 'min_samples_split': 5, 'n_estimators': 100} - Accuracy: 0.7216107242418389
 Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 2, 'n_estimators': 10} - Accuracy: 0.7206394371588057
 Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 2, 'n_estimators': 50} - Accuracy: 0.7215487583375035
 Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 2, 'n_estimators': 100} - Accuracy: 0.7214363535324007
 Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 5, 'n_estimators': 10} - Accuracy: 0.7207907491573353
 Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 5, 'n_estimators': 50} - Accuracy: 0.7214291481635365
 Parameters: {'max_depth': 10, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 5, 'n_estimators': 100} - Accuracy: 0.7214939973036193
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 1,
 'min_samples_split': 2, 'n_estimators': 10} - Accuracy: 0.7039776693063252
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 1,
 'min_samples_split': 2, 'n_estimators': 50} - Accuracy: 0.7104740574863995
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 1,
 'min_samples_split': 2, 'n_estimators': 100} - Accuracy: 0.711479930720021
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 1,
 'min_samples_split': 5, 'n_estimators': 10} - Accuracy: 0.7074463484748386
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 1,
 'min_samples_split': 5, 'n_estimators': 50} - Accuracy: 0.7128936314261616
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 1,
 'min_samples_split': 5, 'n_estimators': 100} - Accuracy: 0.7137049579560032
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 2, 'n_estimators': 10} - Accuracy: 0.7087188237311077
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 2, 'n_estimators': 50} - Accuracy: 0.7137539555504049
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 2, 'n_estimators': 100} - Accuracy: 0.7141661037438711
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 5, 'n_estimators': 10} - Accuracy: 0.7089955087719887
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 5, 'n_estimators': 50} - Accuracy: 0.7140464950859101
 Parameters: {'max_depth': 20, 'max_features': 'auto', 'min_samples_leaf': 2,
 'min_samples_split': 5, 'n_estimators': 100} - Accuracy: 0.7147526248792502

```
[23]: from sklearn.tree import DecisionTreeClassifier

dt = DecisionTreeClassifier()

dt.fit(X_train_scaled, y_train)
y_pred_dt = dt.predict(X_test_scaled)

accuracy_dt = accuracy_score(y_test, y_pred_dt)
print("Initial Decision Tree Accuracy:", accuracy_dt)
```

Initial Decision Tree Accuracy: 0.637678381686371

```
[24]: from sklearn.model_selection import GridSearchCV

param_grid_dt = {
    'max_depth': [None, 10, 20, 30, 40],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': [None, 'auto', 'sqrt', 'log2']
}

grid_search_dt = GridSearchCV(estimator=dt, param_grid=param_grid_dt, cv=5,
    ↪n_jobs=-1, verbose=2)
grid_search_dt.fit(X_train_scaled, y_train)

print("Best Parameters for Decision Tree:", grid_search_dt.best_params_)

best_dt = grid_search_dt.best_estimator_
y_pred_best_dt = best_dt.predict(X_test_scaled)

best_accuracy_dt = accuracy_score(y_test, y_pred_best_dt)
print("Best Accuracy for Decision Tree:", best_accuracy_dt)
```

Fitting 5 folds for each of 180 candidates, totalling 900 fits

Best Parameters for Decision Tree: {'max_depth': 10, 'max_features': None, 'min_samples_leaf': 2, 'min_samples_split': 2}

Best Accuracy for Decision Tree: 0.7201643599779418

```
[28]: print("Best Accuracy for Logistic Regression:", sorted_results[:1])
print("Best Accuracy for Random Forest:", best_accuracy)
print("Best Accuracy for Decision Tree:", best_accuracy_dt)

model_accuracies = {
    "Logistic Regression": sorted_results[:1],
    "Random Forest": best_accuracy,
    "Decision Tree": best_accuracy_dt
}
```

```
sorted_accuracies = sorted(model_accuracies.items(), key=lambda x: x[1],
                             ↪reverse=True)
for model, accuracy in sorted_accuracies:
    print(f"Model: {model} - Accuracy: {accuracy}")
```

Best Accuracy for Logistic Regression: [{'C': 1, 'solver': 'sag', 'max_iter': 200, 'accuracy': 0.7129618421229607}]
 Best Accuracy for Random Forest: 0.721761556981264
 Best Accuracy for Decision Tree: 0.7201643599779418

```
[37]: print(original_model_accuracies)
accuracy_lr = original_model_accuracies['Logistic Regression']
accuracy_rf = original_model_accuracies['Random Forest']
accuracy_knn = original_model_accuracies['K-Nearest Neighbors']

print("Accuracy for Logistic Regression:", accuracy_lr)
print("Accuracy for Random Forest:", accuracy_rf)
print("Accuracy for K-Nearest Neighbors:", accuracy_knn)
```

{'Logistic Regression': 0.7129416670029187, 'Random Forest': 0.6786607755316144, 'K-Nearest Neighbors': 0.682315834779217}
 Accuracy for Logistic Regression: 0.7129416670029187
 Accuracy for Random Forest: 0.6786607755316144
 Accuracy for K-Nearest Neighbors: 0.682315834779217

```
[39]: import matplotlib.pyplot as plt
import numpy as np

original_accuracy_lr = accuracy_lr
original_accuracy_rf = accuracy_rf
original_accuracy_knn = accuracy_knn

best_accuracy_lr = sorted_results[0]['accuracy']
best_accuracy_rf = best_accuracy

original_model_accuracies = {
    "Logistic Regression": original_accuracy_lr,
    "Random Forest": original_accuracy_rf,
    "Decision Tree": accuracy_dt
}

hyper_tuned_accuracies = {
    "Logistic Regression": best_accuracy_lr,
    "Random Forest": best_accuracy_rf,
    "Decision Tree": best_accuracy_dt
}
```

```

labels = original_model_accuracies.keys()

original_scores = original_model_accuracies.values()
tuned_scores = hyper_tuned_accuracies.values()

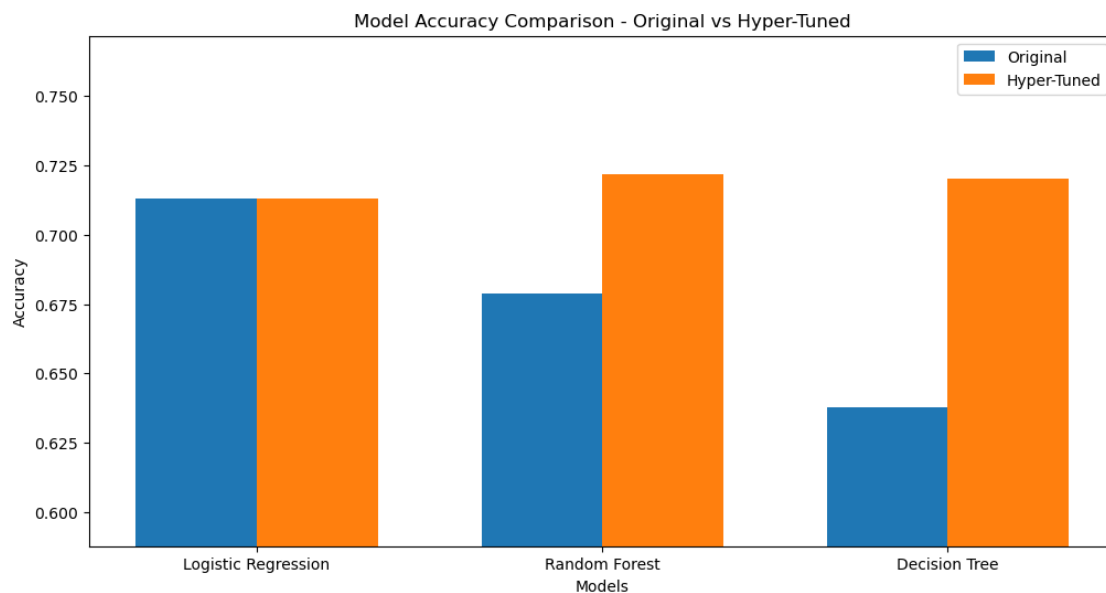
x = np.arange(len(labels))
width = 0.35

plt.figure(figsize=(12, 6))
plt.bar(x - width/2, original_scores, width, label='Original')
plt.bar(x + width/2, tuned_scores, width, label='Hyper-Tuned')

plt.xlabel('Models')
plt.ylabel('Accuracy')
plt.title('Model Accuracy Comparison - Original vs Hyper-Tuned')
plt.xticks(x, labels)
plt.ylim([min(min(original_scores), min(tuned_scores)) - 0.05,
          max(max(original_scores), max(tuned_scores)) + 0.05])
plt.legend()

plt.show()

```



[]: